27 October 2026

Dear Hiring team,

I’m applying for the Systems Development Engineer role on AWS Incident Response in Dublin. Your team’s aim of making each event shorter, or preventing the next one entirely, describes how I already work. At Reward I’m the go-to engineer for production incidents on banking systems running on AWS and Kubernetes, and I measure a fix by whether the problem comes back.

Some examples:

- **Triage to root cause:** I traced recurring daily out-of-memory crashes in a high-volume ingestion service to cache loading, and changed it to run once per batch instead of once per file.
- **Fix and harden:** after resolving a customer-facing outage caused by malformed load-balancer naming in our Kubernetes/Istio ingress, I corrected the Helm templates so it couldn’t recur.
- **Make the next one shorter:** I write runbooks and SOPs so the production services team can resolve repeat issues without engineering, and I’ve automated manual scheduled jobs with EventBridge, Lambda and Fargate.

I also like that AIR engineers build the tooling they use. In my own time I’m building Incident Copilot, which takes an incident description, pulls the matching runbook content and suggests likely root causes and next steps.

I’m AWS Certified Developer – Associate, and I’d be glad to talk about how I could help AIR.

Kind regards,

Padraig Lavery