27 October 2026

Dear Hiring team,

I’m applying for the Forward Deployed Engineer role. Getting LLM-powered agents working inside real client environments needs two things: hands-on AI skills, and the production discipline to keep a system running once it’s live. I’ve spent four years on the second, and I’m building the first.

At Reward I build and run systems for UK and Middle East banking clients on AWS and Kubernetes, each with its own deployment:

- I planned testing for AWS SDK v3 and Node 22 runtime upgrades across repositories and multiple bank deployments, and wrote the engineering migration guide.
- I designed audit logging of staff access to customer accounts across bank portals to meet each client’s security audit requirements.
- I’m the go-to engineer for production incidents: I fixed a Kubernetes/Istio ingress outage and stopped recurring memory crashes in a high-volume ingestion service.

On the AI side, I’m building Incident Copilot, an agentic triage service designed to retrieve runbook content with pgvector, suggest likely root causes, and use tool-calling to check live service status. The knowledge base and API are in place; retrieval and model integration are in progress. I moved it from AWS Bedrock to local models via Ollama to remove inference costs. I’m also preparing for the AWS Certified AI Practitioner exam in November.

I’m based in Belfast, and I’d like to confirm that Northern Ireland meets the “based in Ireland” requirement. Thank you for considering my application.

Kind regards,

Padraig Lavery